

Synchronizing Threads with POSIX Semaphores

1. Why semaphores?
2. Posix semaphores are easy to use
 - [sem_init](#)
 - [sem_wait](#)
 - [sem_post](#)
 - [sem_getvalue](#)
 - [sem_destroy](#)
3. [Exercises](#) [1](#) [2](#)

Now it is time to take a look at some code that does something a little unexpected. The program [badcnt.c](#) creates two new threads, both of which increment a global variable called `cnt` exactly NITER, with NITER = 1,000,000. But the program produces unexpected results.

Exercise 1. Create a directory called `posixsem` in your class Unix directory. Download in this directory the code [badcnt.c](#) and compile it using

```
gcc badcnt.c -o badcnt -lpthread
```

Run the executable `badcnt` and observe the output. Try it on both `tanner` and `felix`.

Quite unexpected! Since `cnt` starts at 0, and both threads increment it NITER times, we should see `cnt` equal to $2 \times \text{NITER}$ at the end of the program. What happens?

Threads can greatly simplify writing elegant and efficient programs. However, there are problems when multiple threads share a common address space, like the variable `cnt` in our earlier example.

To understand what might happen, let us analyze this simple piece of code:

THREAD 1	THREAD 2
<code>a = data;</code>	<code>b = data;</code>
<code>a++;</code>	<code>b--;</code>
<code>data = a;</code>	<code>data = b;</code>

Now if this code is executed serially (for instance, THREAD 1 first and then THREAD 2), there are no problems. However threads execute in an arbitrary order, so consider the following situation:

Thread 1	Thread 2	data
<code>a = data;</code>	---	0
<code>a = a+1;</code>	---	0
---	<code>b = data; // 0</code>	0
---	<code>b = b + 1;</code>	0
<code>data = a; // 1</code>	---	1
---	<code>data = b; // 1</code>	1

So data could end up +1, 0, -1, and there is **NO WAY** to know which value! It is completely non-deterministic!

The solution to this is to provide functions that will block a thread if another thread is accessing data that it is using.

Pthreads may use semaphores to achieve this.

Posix semaphores

All POSIX semaphore functions and types are prototyped or defined in [semaphore.h](#). To define a semaphore object, use

```
sem_t sem_name;
```

To initialize a semaphore, use `sem_init()`:

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

- `sem` points to a semaphore object to initialize
- `pshared` is a flag indicating whether or not the semaphore should be shared with `fork()`ed processes. LinuxThreads does not currently support shared semaphores
- `value` is an initial value to set the semaphore to

Example of use:

```
sem_init(&sem_name, 0, 10);
```

To wait on a semaphore, use `sem_wait()`:

```
int sem_wait(sem_t *sem);
```

Example of use:

```
sem_wait(&sem_name);
```

- If the value of the semaphore is negative, the calling process blocks; one of the blocked processes wakes up when another process calls `sem_post`.
-

To increment the value of a semaphore, use `sem_post()`:

```
int sem_post(sem_t *sem);
```

Example of use:

```
sem_post(&sem_name);
```

- It increments the value of the semaphore and wakes up a blocked process waiting on the semaphore, if any.
-

To find out the value of a semaphore, use

```
int sem_getvalue(sem_t *sem, int *valp);
```

- gets the current value of `sem` and places it in the location pointed to by `valp`

Example of use:

```
int value;  
sem_getvalue(&sem_name, &value);  
printf("The value of the semaphors is %d\n", value);
```

To destroy a semaphore, use

```
int sem_destroy(sem_t *sem);
```

- destroys the semaphore; no threads should be waiting on the semaphore if its destruction is to succeed.

Example of use:

```
sem_destroy(&sem_name);
```

Using semaphores - a short example

Consider the problem we had before and now let us use semaphores:

Declare the semaphore global (outside of any function):

```
sem_t mutex;
```

Initialize the semaphore in the main function:

```
sem_init(&mutex, 0, 1);
```

Thread 1	Thread 2	data
sem_wait (&mutex);	---	0
---	sem_wait (&mutex);	0
a = data;	/* blocked */	0
a = a+1;	/* blocked */	0
data = a;	/* blocked */	1
sem_post (&mutex);	/* blocked */	1
/* blocked */	b = data;	1
/* blocked */	b = b + 1;	1
/* blocked */	data = b;	2
/* blocked */	sem_post (&mutex);	2
[data is fine. The data race is gone.]		

Exercise 2. Use the example above as a guide to fix the program [badcnt.c](#), so that the program always produces the expected output (the value 2*NITER). Make a copy of badcnt.c into goodcnt.c before you modify the code.

To compile a program that uses pthreads *and* posix semaphores, use

```
gcc -o filename filename.c -lpthread -lrt
```
